

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

In [2]: df = pd.read_csv('Students.csv')
df
```

	Background	Machine Learning	Python	Remark
0	Tech	78.0	98.0	Pass
1	Tech	NaN	10.0	Fail
2	Non-Tech	69.0	NaN	Pass
3	Non-Tech	76.0	56.0	Pass
4	Tech	NaN	NaN	Pass
5	Tech	20.0	40.0	Fail
6	Tech	NaN	10.0	Fail
7	Non-Tech	54.0	20.0	Pass
8	Tech	87.0	98.0	Pass
9	Non-Tech	NaN	NaN	Fail

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Background       10 non-null     object
1   Machine Learning  6 non-null      float64
2   Python           7 non-null      float64
3   Remark           10 non-null     object
dtypes: float64(2), object(2)
memory usage: 448.0+ bytes

In [4]: x = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

```
In [5]: df

Out[5]:
```

	Background	Machine Learning	Python	Remark
0	Tech	78.0	98.0	Pass
1	Tech	NaN	10.0	Fail
2	Non-Tech	69.0	NaN	Pass
3	Non-Tech	76.0	56.0	Pass
4	Tech	NaN	NaN	Pass
5	Tech	20.0	40.0	Fail
6	Tech	NaN	10.0	Fail
7	Non-Tech	54.0	20.0	Pass
8	Tech	87.0	98.0	Pass
9	Non-Tech	NaN	NaN	Fail

```
In [6]: from sklearn.impute import SimpleImputer
si = SimpleImputer(missing_values=np.nan, strategy="mean")
df.iloc[:, 1:3]=si.fit_transform(df.iloc[:, 1:3])

In [7]: df
```

	Background	Machine Learning	Python	Remark
0	Tech	78.0	98.000000	Pass
1	Tech	64.0	10.000000	Fail
2	Non-Tech	69.0	47.428571	Pass
3	Non-Tech	76.0	56.000000	Pass
4	Tech	64.0	47.428571	Pass
5	Tech	20.0	40.000000	Fail
6	Tech	64.0	10.000000	Fail
7	Non-Tech	54.0	20.000000	Pass
8	Tech	87.0	98.000000	Pass
9	Non-Tech	64.0	47.428571	Fail

- Use OneHotEncoder to Features

```
In [8]: background = pd.get_dummies(x['Background'])

In [9]: df = pd.concat([background, df], axis=1)

In [10]: df
```

	Non-Tech	Tech	Background	Machine Learning	Python	Remark
0	0	1	Tech	78.0	98.000000	Pass
1	0	1	Tech	64.0	10.000000	Fail
2	1	0	Non-Tech	69.0	47.428571	Pass
3	1	0	Non-Tech	76.0	56.000000	Pass
4	0	1	Tech	64.0	47.428571	Pass
5	0	1	Tech	20.0	40.000000	Fail
6	0	1	Tech	64.0	10.000000	Fail
7	1	0	Non-Tech	54.0	20.000000	Pass
8	0	1	Tech	87.0	98.000000	Pass
9	1	0	Non-Tech	64.0	47.428571	Fail

```
In [11]: df.drop('Background', axis=1, inplace=True)

In [12]: df
```

	Non-Tech	Tech	Machine Learning	Python	Remark
0	0	1	78.0	98.000000	Pass
1	0	1	64.0	10.000000	Fail
2	1	0	69.0	47.428571	Pass
3	1	0	76.0	56.000000	Pass
4	0	1	64.0	47.428571	Pass
5	0	1	20.0	40.000000	Fail
6	0	1	64.0	10.000000	Fail
7	1	0	54.0	20.000000	Pass
8	0	1	87.0	98.000000	Pass
9	1	0	64.0	47.428571	Fail

- Use LabelEncoder to Target

```
In [13]: y = df[["Remark"]]

In [14]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y[["Remarks"]] = le.fit_transform(y[["Remark"]])

In [15]: df = pd.concat([df, y], axis=1)

In [16]: df

Out[16]:
```

	Non-Tech	Tech	Machine Learning	Python	Remark	Remark	Remarks
0	0	1	78.0	98.000000	Pass	Pass	1
1	0	1	64.0	10.000000	Fail	Fail	0
2	1	0	69.0	47.428571	Pass	Pass	1
3	1	0	76.0	56.000000	Pass	Pass	1
4	0	1	64.0	47.428571	Pass	Pass	1
5	0	1	20.0	40.000000	Fail	Fail	0
6	0	1	64.0	10.000000	Fail	Fail	0
7	1	0	54.0	20.000000	Pass	Pass	1
8	0	1	87.0	98.000000	Pass	Pass	1
9	1	0	64.0	47.428571	Fail	Fail	0

```
In [17]: df.drop('Remark', axis=1, inplace=True)

In [18]: df
```

	Non-Tech	Tech	Machine Learning	Python	Remarks
0	0	1	78.0	98.000000	1
1	0	1	64.0	10.000000	0
2	1	0	69.0	47.428571	1
3	1	0	76.0	56.000000	1
4	0	1	64.0	47.428571	1
5	0	1	20.0	40.000000	0
6	0	1	64.0	10.000000	0
7	1	0	54.0	20.000000	1
8	0	1	87.0	98.000000	1
9	1	0	64.0	47.428571	0

Do Feature Scaling with StandardScaler

```
In [19]: from sklearn.preprocessing import StandardScaler
for col in df:
    ss = StandardScaler()
    df[col]=ss.fit_transform(df[[col]])

In [20]: df

Out[20]:
```

	Non-Tech	Tech	Machine Learning	Python	Remarks
0	-0.816497	0.816497	0.817889	1.704984	0.816497
1	-0.816497	0.816497	0.000000	-1.261881	-1.224745
2	1.224745	-1.224745	0.292103	0.000000	0.816497
3	1.224745	-1.224745	0.701047	0.288980	0.816497
4	-0.816497	0.816497	0.000000	0.000000	0.816497
5	-0.816497	0.816497	-2.570507	-0.250450	-1.224745
6	-0.816497	0.816497	0.000000	-1.261881	-1.224745
7	1.224745	-1.224745	-0.584206	-0.924737	0.816497
8	-0.816497	0.816497	1.343674	1.704984	0.816497
9	1.224745	-1.224745	0.000000	0.000000	-1.224745

```
In [21]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Non-Tech        10 non-null     float64
1   Tech            10 non-null     float64
2   Machine Learning 10 non-null     float64
3   Python          10 non-null     float64
4   Remarks         10 non-null     float64
dtypes: float64(5)
memory usage: 528.0 bytes

In [25]: x = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

- Split the Training and Testing set

```
In [26]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.3, random_state=1)

In [27]: from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(xtrain, ytrain)
ypred = linreg.predict(xtest)
```

Checking Accuracy

```
In [28]: from sklearn.metrics import mean_absolute_error as mae , mean_squared_error as mse , r2_score

In [29]: print(f"MAE :- {mae(ytest, ypred)}")
print(f"MSE :- {mse(ytest, ypred)}")
print(f"RMSE :- {mse(ytest, ypred)**0.5}")
print(f"Accuracy :- {r2_score(ytest, ypred)}")

MAE :- 0.9679826589808663
MSE :- 1.689331803567157
RMSE :- 1.29948598186753
Accuracy :- -0.8026616747852522

In [ ] :
```